


COMP9121: Design of Networks and Distributed Systems

Week 5: Network Layer 2
Wei Bao
School of Computer Science

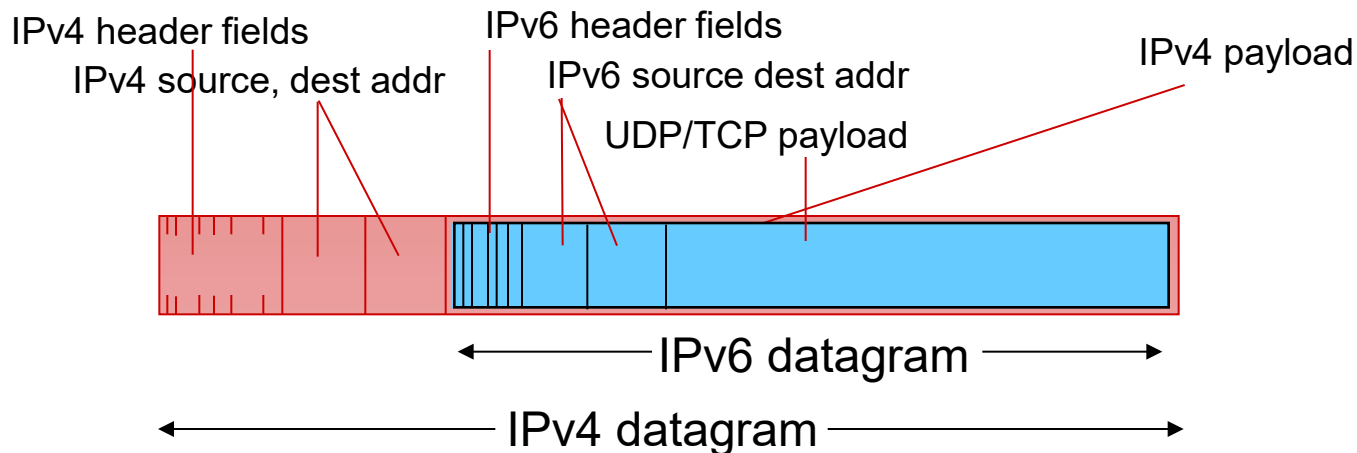


THE UNIVERSITY OF
SYDNEY




Transition from IPv4 to IPv6

- not all routers can be upgraded simultaneously
 - no “flag days”
 - how will network operate with mixed IPv4 and IPv6 routers?
- **tunneling**: IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers




Tunneling

logical view:




physical view:

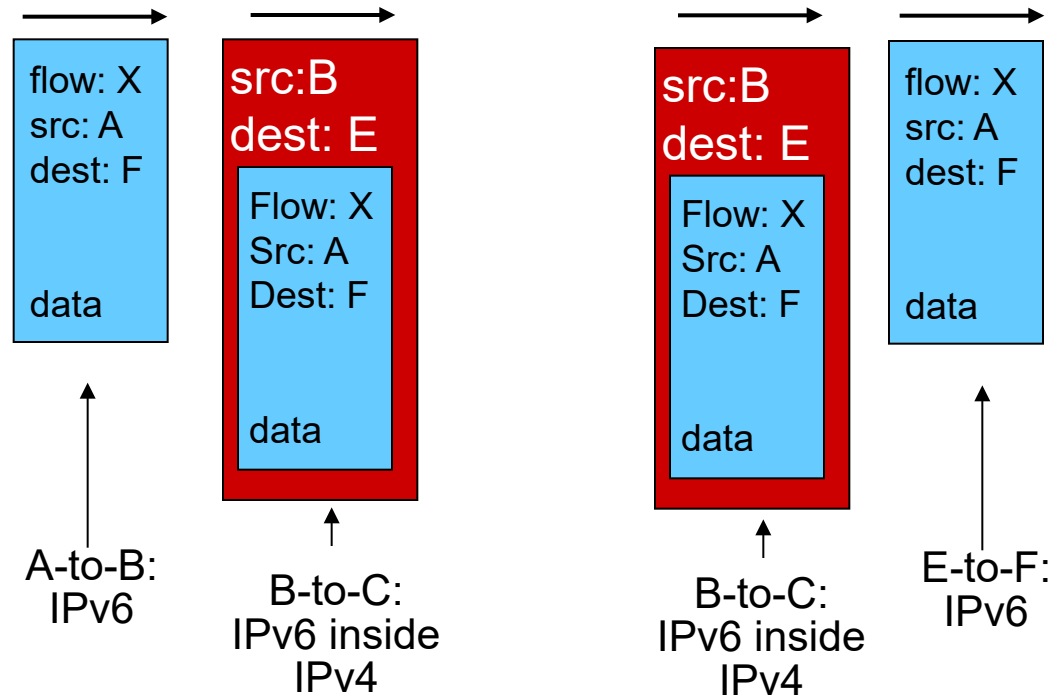



Tunneling

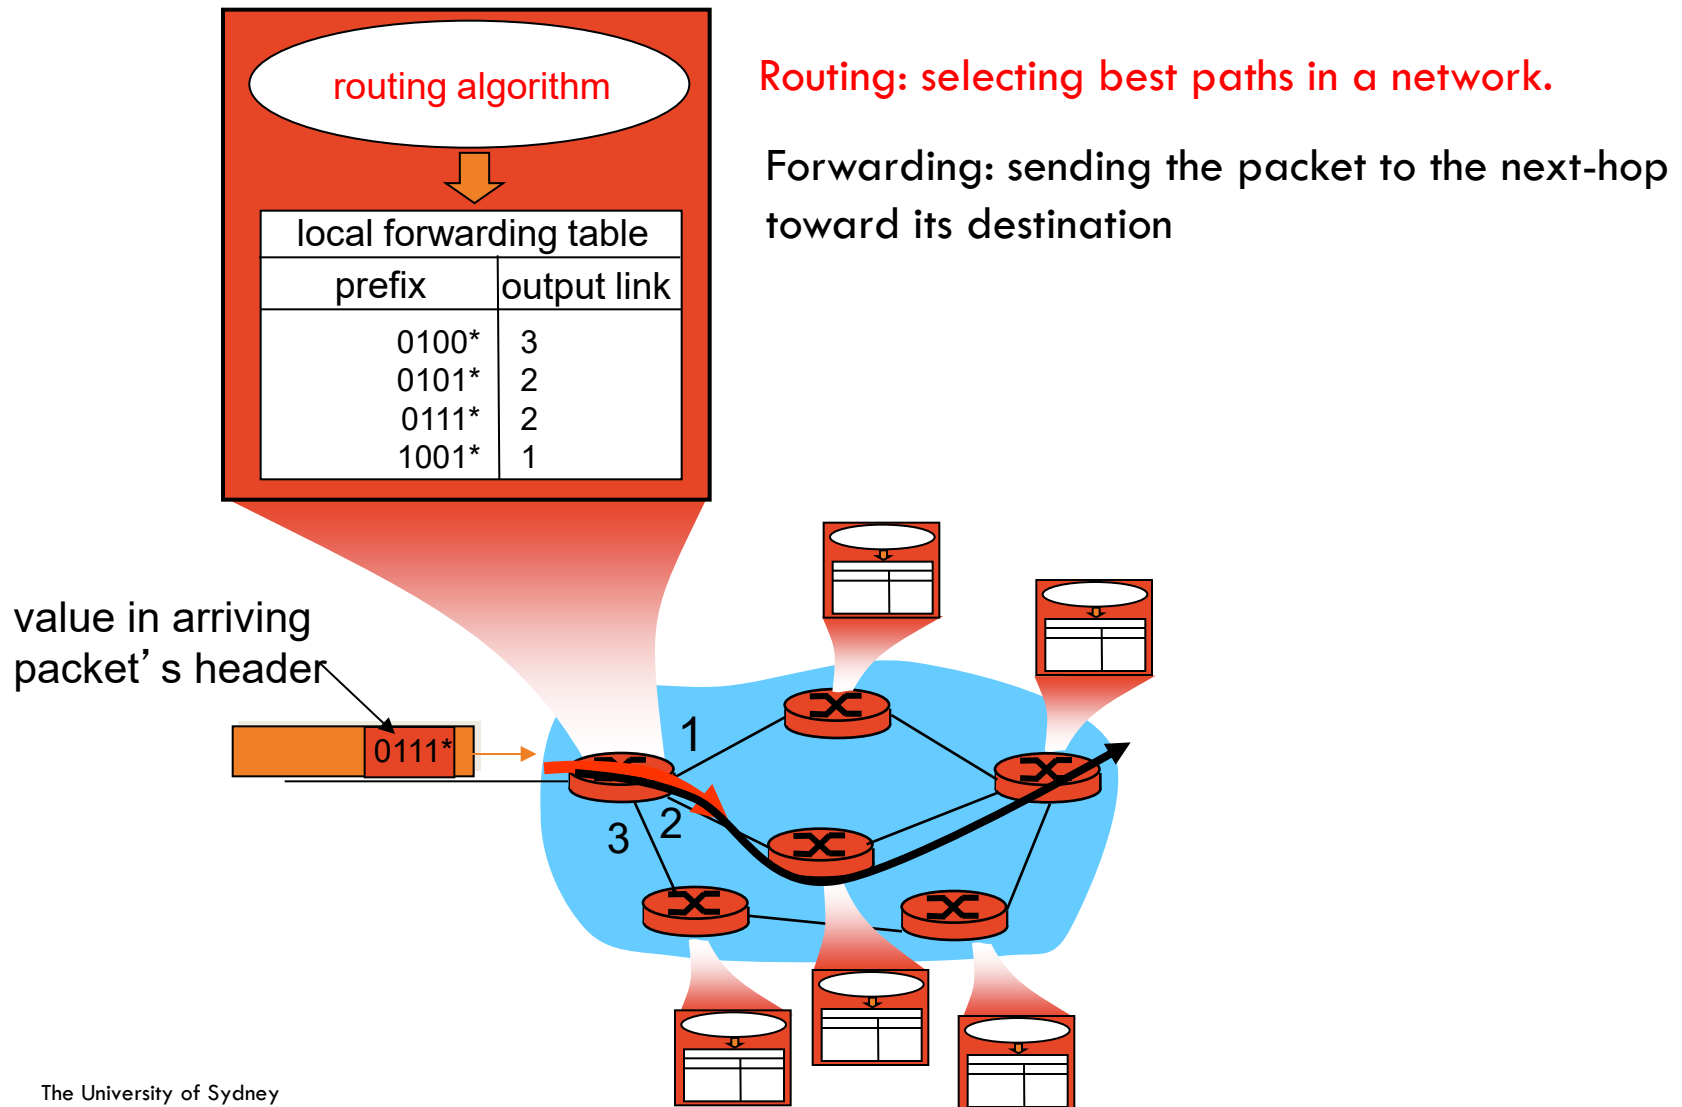
logical view:



physical view:



Routing and forwarding




Link Cost

i : i -th node (router).

C_{ij} : link cost between i and j

$C_{ij} = \text{infinity}$ if i and j are not connected



Aim: Shortest path.

Classification

Global or decentralized information?

Global:

- all routers have complete topology, link cost info
- “link state” algorithms

Decentralized:

- router only knows physically-connected neighbors, link costs to neighbors
- “distance vector” algorithms

Link State Algorithm: Dijkstra

Link-State Algorithm


- Basic idea: **two step** procedure
 - 1 Each source node gets a **map** of all nodes and link costs of the entire network
 - 2 Find the **shortest path** on the map from the source node to all destination nodes
- Step 1: Broadcast of link-state information
 - Every node i in the network broadcasts to every other node in the network:
 - ID's of its neighbors: \mathcal{N}_i = set of neighbors of i
 - Distances to its neighbors: $\{C_{ij} \mid i \in \mathcal{N}_i\}$
 - **Flooding** is a popular method of broadcasting packets
- Step 2: Dijkstra

Dijkstra Algorithm: Finding shortest paths in order

Find shortest paths from source s to all other destinations

Closest node to s is 1 hop away
2nd closest node to s is 1 hop away from s or w''


Otherwise: suppose a is the 2nd closest node: Contradiction!



Dijkstra Algorithm: Finding shortest paths in order


Find shortest paths from source s to all other destinations

Closest node to s is 1 hop away
2nd closest node to s is 1 hop away from s or w''
3rd closest node to s is 1 hop away from s , w'' , or x




Dijkstra's algorithm

- C_{ij} : distance/cost from i to j
- D_i : current shortest distance from s to i
- N : set of nodes for which shortest path already found
- **Initialization: (Start with source node s)**
 - $N = \{s\}$, $D_s = 0$, “ s is distance zero from itself”
 - $D_i = C_{si}$ for all $i \neq s$, distances of directly-connected neighbors




Dijkstra's algorithm

- **Step A: (Find next closest node i)**
 - Find $i \notin N$ such that
 - $D_i = \min D_j$ for $j \notin N$
 - Add i to N
 - If N contains all the nodes, stop



Dijkstra's algorithm


- (node i was added to N last step)
 - Other nodes may find a better way via the newly added
 - **Step B: (update minimum costs)**
 - For each node $j \notin N$
 - $D_j = \min (D_j, D_i + C_{ij})$
- Minimum distance from s to j through node i in N*



Dijkstra's algorithm

- N : set of nodes for which shortest path already found
- **Initialization: (Start with source node s)**
 - $N = \{s\}$, $D_s = 0$, “ s is distance zero from itself”
 - $D_i = C_{si}$ for all $i \neq s$, distances of directly-connected neighbors
- **Step A: (Find next closest node i)**
 - Find $i \notin N$ such that
 - $D_i = \min D_j$ for $j \notin N$
 - Add i to N
 - If N contains all the nodes, stop
- **Step B: (update minimum costs)**
 - For each node $j \notin N$
 - $D_j = \min (D_j, D_i + C_{ij})$
 - Go to Step A


Dijkstra's algorithm: Example



Find the shortest path from 1 to all other nodes.

Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial						
1						
2						
3						
4						
5						


Dijkstra's algorithm: Example



(next node/last node, cost D_j)

Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	(-1, ∞)	(-1, ∞)
1						
2						
3						
4						
5						


Dijkstra's algorithm: Example



(next node/last node, cost)


Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	(-1, ∞)	(-1, ∞)
1	{1, 3}					
2						
3						
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2						
3						
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}					
3						
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}			(3,4)	(2,7)	(3,3)
3						
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}			(3,4)	(2,7)	(3,3)
3	{1,2,3,6}					
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}			(3,4)	(2,7)	(3,3)
3	{1,2,3,6}			(3,4)	(6,5)	
4						
5						

Dijkstra's algorithm: Example




Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}			(3,4)	(2,7)	(3,3)
3	{1,2,3,6}			(3,4)	(6,5)	
4	{1,2,3,4,6}				(6,5)	

Dijkstra's algorithm: Example



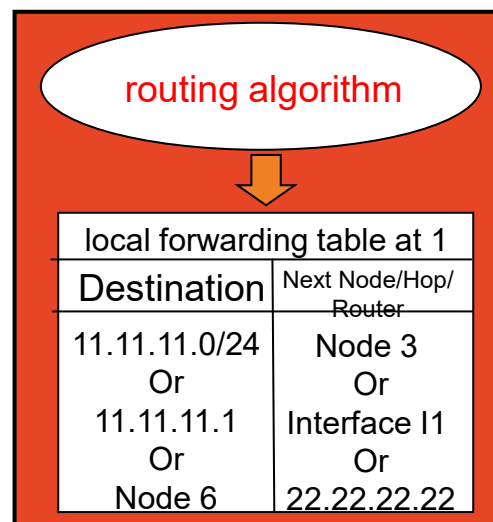
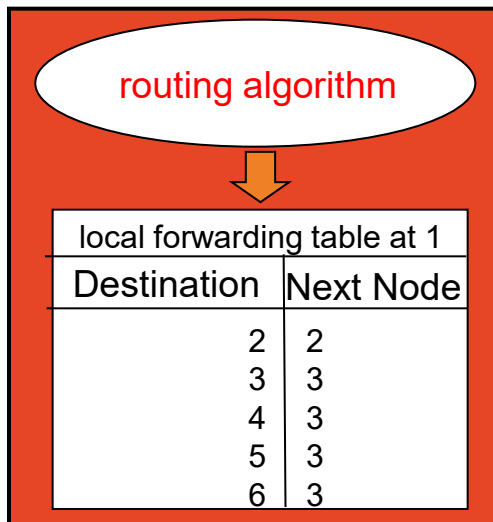
Iteration	Tree	N_2	N_3	N_4	N_5	N_6
Initial	{1}	(1,3)	(1,2)	(1,5)	$(-1, \infty)$	$(-1, \infty)$
1	{1,3}	(1,3)		(3,4)	$(-1, \infty)$	(3,3)
2	{1,2,3}			(3,4)	(2,7)	(3,3)
3	{1,2,3,6}			(3,4)	(6,5)	
4	{1,2,3,4,6}				(6,5)	
5	{1,2,3,4,5,6}	(1,3)	(1,2)	(3,4)	(6,5)	(3,3)

Forwarding table



For presentation convenience
Destination can be: Node name,
IP address, IP prefix, etc.

Destination can be: Node name,
IP address, interface, etc.



Reaction to Failure

- If a link fails,
 - Router sets link distance to infinity & floods the network with an update packet
 - All routers immediately update their link database & recalculate their shortest paths
 - Recovery very quick
- But watch out for **old** update messages
 - Add time stamp or sequence # to each update message
 - Check whether each received update message is new
 - If new, add it to database and broadcast
 - If older, send update message on arriving link

Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- each iteration: need to check all nodes, w , not in N
- operations: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$
 - Using a heap to find the min.


Important Concepts

- Link-state is a two-step procedure
 - Learn the map (by flooding)
 - Find the shortest path (use Dijkstra)
- LS is centralized!

Distance Vector Algorithm


Shortest Path to SJ

Focus on how nodes find their shortest paths to a given destination node, i.e. SJ



But we don't know the shortest paths

i only has local info
from neighbours



Distance Vector Algorithm

Bellman-Ford Equation (dynamic programming)

Define

$D_x(y) :=$ cost of shortest path from x to y

Then

$$D_x(y) = \min_v \{C(x,v) + D_v(y)\}$$

where min is taken over all neighbors v of x

Destination node 6

$$D1 = \min \{ 3 + D2, 2 + D3, 5 + D4 \}$$

$$D2 = \min \{ 3 + D1, 1 + D4, 4 + D5 \}$$


$$D3 = \min \{ 2 + D1, 2 + D4, 1 \}$$

$$D4 = \min \{ 5 + D1, 1 + D2, 2 + D3, 3 + D5 \}$$

$$D5 = \min \{ 4 + D2, 3 + D4, 2 \}$$

How to solve: Use an iterative procedure!

Use this Bellman-ford equation every round, until convergence.




Iteration
Initial
1
2
3

Node 1
$(-1, \infty)$

Node 2
$(-1, \infty)$

Node 3
$(-1, \infty)$

Node 4
$(-1, \infty)$




San Jose

Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$
1					
2					
3					

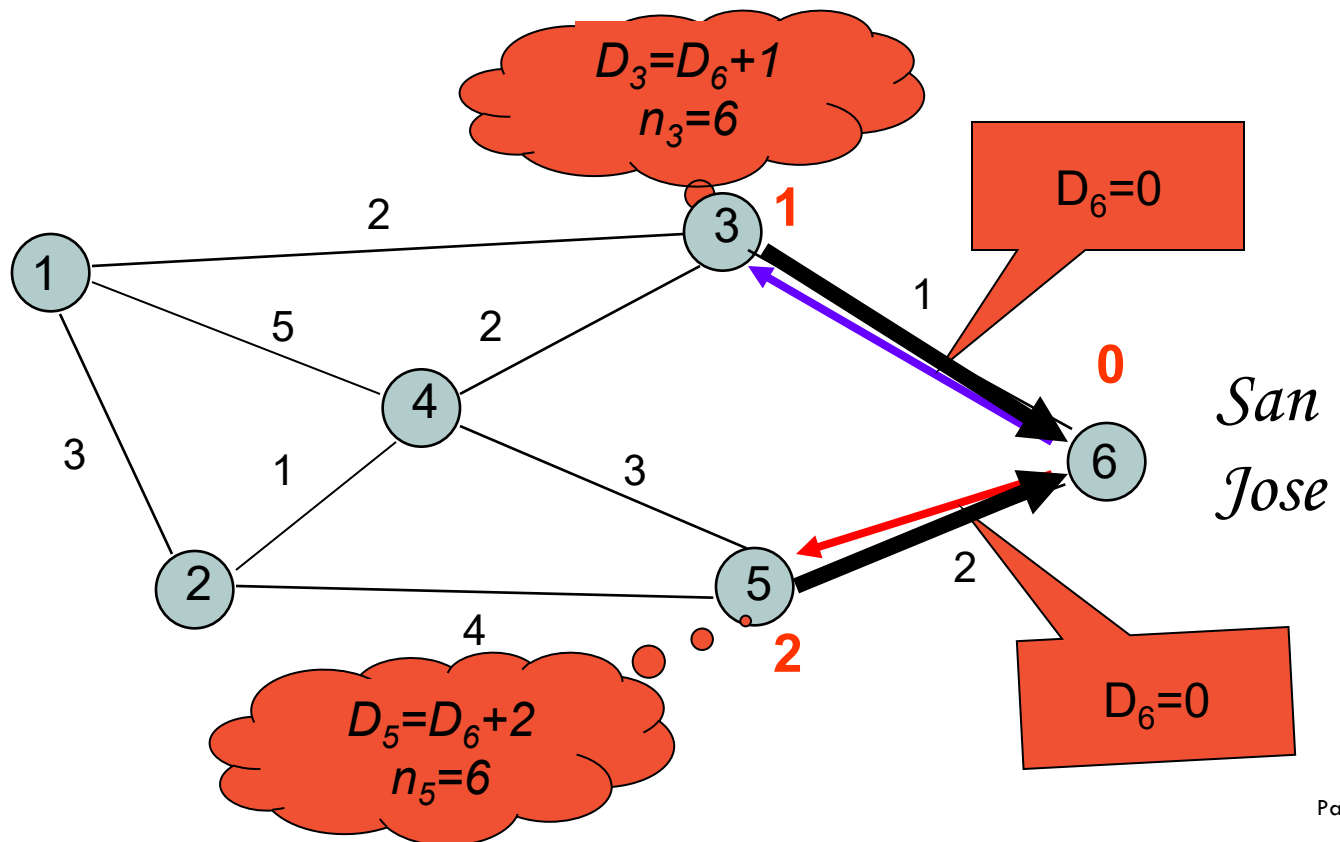
Table entry
@ node 1
for dest SJ

Table entry
@ node 3
for dest SJ




*San
Jose*


Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$
1	$(-1, \infty)$	$(-1, \infty)$	$(6, 1)$	$(-1, \infty)$	$(6, 2)$
2					
3					




Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$
1	$(-1, \infty)$	$(-1, \infty)$	$(6, 1)$	$(-1, \infty)$	$(6, 2)$
2	$(3, 3)$	$(5, 6)$	$(6, 1)$	$(3, 3)$	$(6, 2)$
3					



Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$
1	$(-1, \infty)$	$(-1, \infty)$	$(6, 1)$	$(-1, \infty)$	$(6, 2)$
2	$(3, 3)$	$(5, 6)$	$(6, 1)$	$(3, 3)$	$(6, 2)$
3					



Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$
1	$(-1, \infty)$	$(-1, \infty)$	$(6, 1)$	$(-1, \infty)$	$(6, 2)$
2	$(3, 3)$	$(5, 6)$	$(6, 1)$	$(3, 3)$	$(6, 2)$
3	$(3, 3)$	$(4, 4)$	$(6, 1)$	$(3, 3)$	$(6, 2)$



Distance vector algorithm

Basic idea:

- From time-to-time (ex. 30 sec), each node sends its own distance vector estimate to neighbours
- Distance vector: (Current shortest distance to 1, Current shortest distance to 2, ... Current shortest distance to N)
- When a node x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{C(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

- Under minor, natural conditions, the estimate $D_x(y)$ converges to the actual least cost $d_x(y)$

Example

Consider the following network in which distance vector routing is used. Each node in this network sends its routing table using a vector of size 6 where each entity of the vector represents the cost to the corresponding node, i.e. (Cost-to-Node1, Cost-to-Node2, Cost-to-Node3, Cost-to-Node4, Cost-to-Node5, Cost-to-Node6).

The following cost vectors have just arrived at router 1 from its neighbours:


from 2: (2, 0, 2, 7, 4, 1);

from 3: (3, 2, 0, 3, 1, 4);


from 5: (1, 3, 1, 4, 0, 3).

The link costs between node 1 and nodes 2, 3, and 5 are 2, 3, and 1, respectively. What is the routing table at node 1?

	Cost	Next Node
For Node 2: $\min \{2+0, 3+2, 1+3\} = 2$	2	2
For Node 3: $\min \{2+2, 3+0, 1+1\} = 2$	5	5
For Node 4: $\min \{2+7, 3+3, 1+4\} = 5$	5	5
For Node 5: $\min \{2+4, 3+1, 1+0\} = 1$	5	5
For Node 6: $\min \{2+1, 3+4, 1+3\} = 3$	2	2




Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	(3,3)	(4,4)	(6, 1)	(3,3)	(6,2)
1	(3,3)	(4,4)	(4, 5)	(3,3)	(6,2)
2					
3					




Network disconnected; Loop created between nodes 3 and 4

Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	(3,3)	(4,4)	(6, 1)	(3,3)	(6,2)
1	(3,3)	(4,4)	(4, 5)	(3,3)	(6,2)
2	(3,7)	(4,4)	(4, 5)	(5,5)	(6,2)
3					




Node 4 could have chosen 2 as next node because of tie

Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
Initial	(3,3)	(4,4)	(6, 1)	(3,3)	(6,2)
1	(3,3)	(4,4)	(4, 5)	(3,3)	(6,2)
2	(3,7)	(4,4)	(4, 5)	(5,5)	(6,2)
3	(3,7)	(4,6)	(4, 7)	(5,5)	(6,2)




Node 2 could have chosen 5 as next node because of tie

Iteration	Node 1	Node 2	Node 3	Node 4	Node 5
1	(3,3)	(4,4)	(4, 5)	(3,3)	(6,2)
2	(3,7)	(4,4)	(4, 5)	(2,5)	(6,2)
3	(3,7)	(4,6)	(4, 7)	(5,5)	(6,2)
4	(2,9)	(4,6)	(4, 7)	(5,5)	(6,2)



Node 1 could have chosen 3 as next node because of tie

Counting to Infinity Problem



Nodes believe best path is through each other

(Destination is node 4)


Update	Node 1	Node 2	Node 3
Before break	(2,3)	(3,2)	(4, 1)
After break	(2,3)	(3,2)	(2,3)
1	(2,3)	(3,4)	(2,3)
2	(2,5)	(3,4)	(2,5)
3	(2,5)	(3,6)	(2,5)
4	(2,7)	(3,6)	(2,7)
5	(2,7)	(3,8)	(2,7)
...

Problem: Bad News Travels Slowly

Remedies

- **Split Horizon**
 - Do not report route to a destination to the neighbor from which route was learned
- **Poisoned Reverse**
 - Report route to a destination to the neighbor from which route was learned, but with infinite distance

Split Horizon with Poison Reverse



Nodes believe best path is through each other

Don't learn from 1 Don't learn from 2

Update	Node 1	Node 2	Node 3	
Before break	(2, 3)	(3, 2)	(4, 1)	
After break	(2, 3)	(3, 2)	$(-1, \infty)$	Node 2 advertizes its route to 4 to node 3 as having distance infinity; node 3 finds there is no route to 4
1	(2, 3)	$(-1, \infty)$	$(-1, \infty)$	Node 1 advertizes its route to 4 to node 2 as having distance infinity; node 2 finds there is no route to 4
2	$(-1, \infty)$	$(-1, \infty)$	$(-1, \infty)$	Node 1 finds there is no route to 4

Important Concepts

- ❑ *Link State* is centralized
 - ❑ Dijkstra
- ❑ *Distance Vector* is decentralized
 - ❑ Split Horizon/Poison Reverse